

Projecting Professional Placement of Summer League Players

Patrick Mellady

Objective

This report is concerned with the answer to the following question:

Given season level data, we would like you to predict the professional outcomes of college players

What follows is the solution to this question for certain players and how this information can be used for the purposes of player recruitment.

Executive Summary

This report evaluates the ability of statistical models to predict the professional placement of NCAA baseball players using college career data. Two approaches were compared: a proportional odds model and random forests.

Key findings include:

- **Counting-based statistics and age** in a player's final season are the strongest predictors of professional placement.
- **Random forests achieved higher accuracy** (approx. 99%) than proportional odds models (approx. 90%), suggesting they capture complex relationships more effectively. However, such a near perfect fit also indicates overfitting and a less generalizable model.
- **Proportional odds models remain valuable** because they quantify the influence of individual statistics, offering interpretability that random forests lack.
- **Scouting Insights** Both models agree on the inherent value certain statistics provide in projecting professional placement. These statistics can be used as a screening tool to find valuable players and utilize scouting resources more efficiently.

These results demonstrate that statistical modeling can complement traditional scouting by identifying which players are most likely to succeed at the professional level and clarifying the attributes that drive those outcomes.

Background

To answer this question, it is necessary to draw distinctions between potential professional outcomes for players. For the purposes of this analysis, professional outcomes for baseball players will have three categories: None, Minors, and Majors. These three categories indicate if a player has no professional placement, is a Minor League player, or is a Major League player, respectively.

Since the predictions will be based on statistical data, and pitchers and hitters have different recorded statistics, we must fit separate models for each of the two types of players. We will use the familiar recorded statistics, a portion of which are shown for hitters and pitchers below.

Some Familiar Batting Statistics

Outcome	Name	Age	PA	AB	R	H	X2B	X3B	HR
None	Joyce, Dylan	22	63	483	416	102	143	20	0

Some Familiar Pitching Statistics

Outcome	Name	Age	W	L	W.L.	ERA	RA9	G	GS
None	Spencer, Casey	19	0	0	NA	8.06	8.06	13	3

Using the above statistics, predictions, which denote projected professional outcomes, will be fit for all players in the data set. The data have been scraped from Baseball-Reference.com and contain College career statistics for 651 hitters and 591 pitchers.

Methodology

This problem can be thought of in terms of a multi-category ordered classification (i.e. the “None” category is a lower order than the “Minors” category, which is a lower order than the “Majors” category and the goals is to find into which category a player will fall). Classically, statisticians have developed methods for analyzing this kind of data by using what is called a *Proportional Odds* model. Intuitively, this works by using each player’s career statistics to estimate the probability of ending up in each of the three categories, and then make the final projection by whichever category has the highest probability. The benefits of the proportional odds model are that it can numerically quantify how influential each predictor is on professional placement and, since we have these quantifications, we can perform statistical tests to determine if certain information is influential at all.

An alternative approach to this type of problem is to use classification trees or random forests. These methods have the benefit of both being easier to understand and not relying on assumptions about the structure of the data and can thus better model complex relationships. However, with these tree based methods, there is no ability to quantify exactly how influential each predictor is on professional placement, which is often desirable.

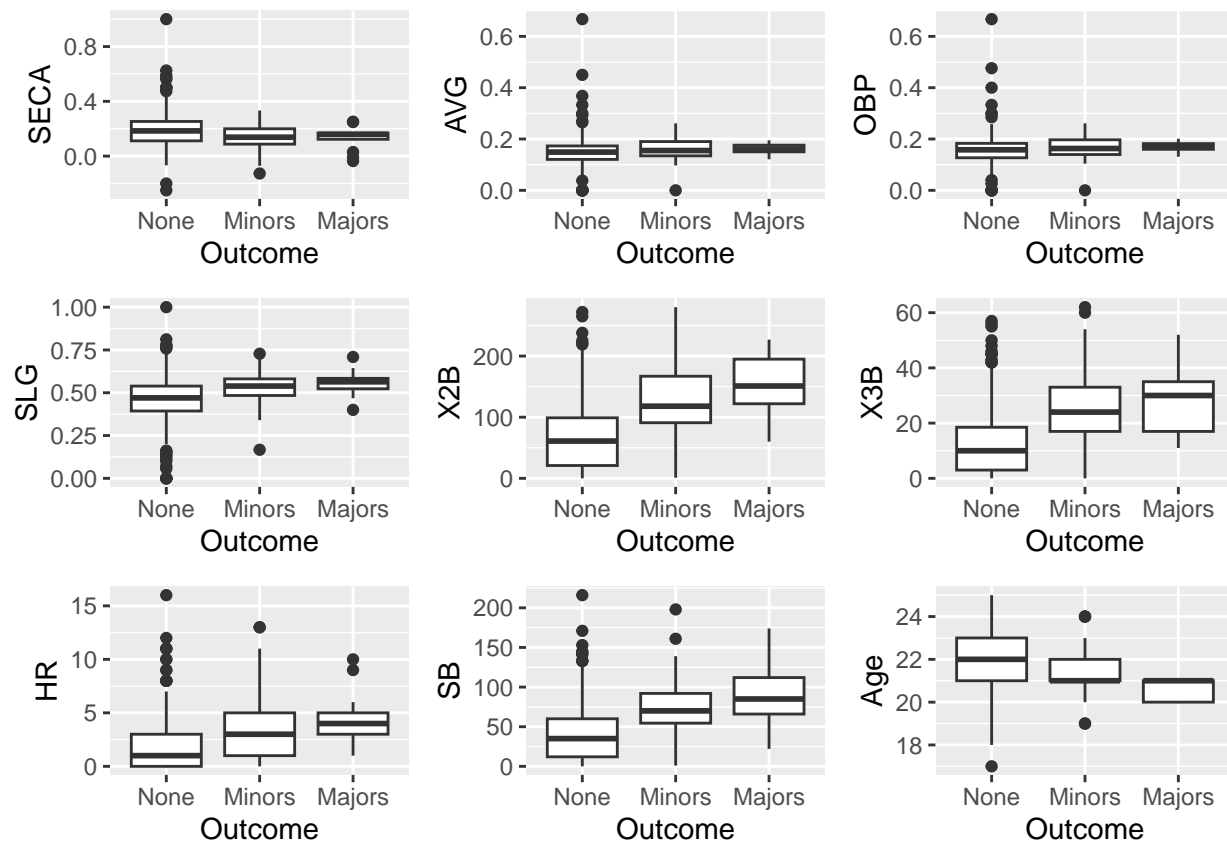
In this report, both the proportional odds and random forests models are fit to the data.

Results

Batters

Exploratory Analysis

To see which statistics are the most related to a batter’s professional outcome, below shows some exploratory data analysis via data visualization.



The above suggests that player placement is more affected by counting-based statistics than rate-based statistics, which may be because of limitations on playing time for some players leading to outliers. Additionally, see that being 21-years-old in your final NCAA season is associated with playing at the professional level, while being older or younger is associated with not being drafted. With these insights, counting-based statistics will be used when fitting the parametric model.

The Proportional Odds Model

The proportional odds model has been fit to the batters data, a random sample of players, their current placements, and their predicted placements is shown below.

Players Actual vs Predicted Placements (PO Model)

Name	Current.Placement	Projected.Placement
Saum, Charlie	None	None
Legg, Jimmy	None	Minors
Atwood, Andy	Minors	Minors
Suddleson, Jake	Minors	Majors
Spohn, Harrison	Minors	Minors
Oakley, Nick	None	None
Oyama, Joichiro	None	None

The following table summarizes the accuracy of this method. This table shows what percent of players from each of the observed outcomes were predicted in each category. For example, of the players who have an

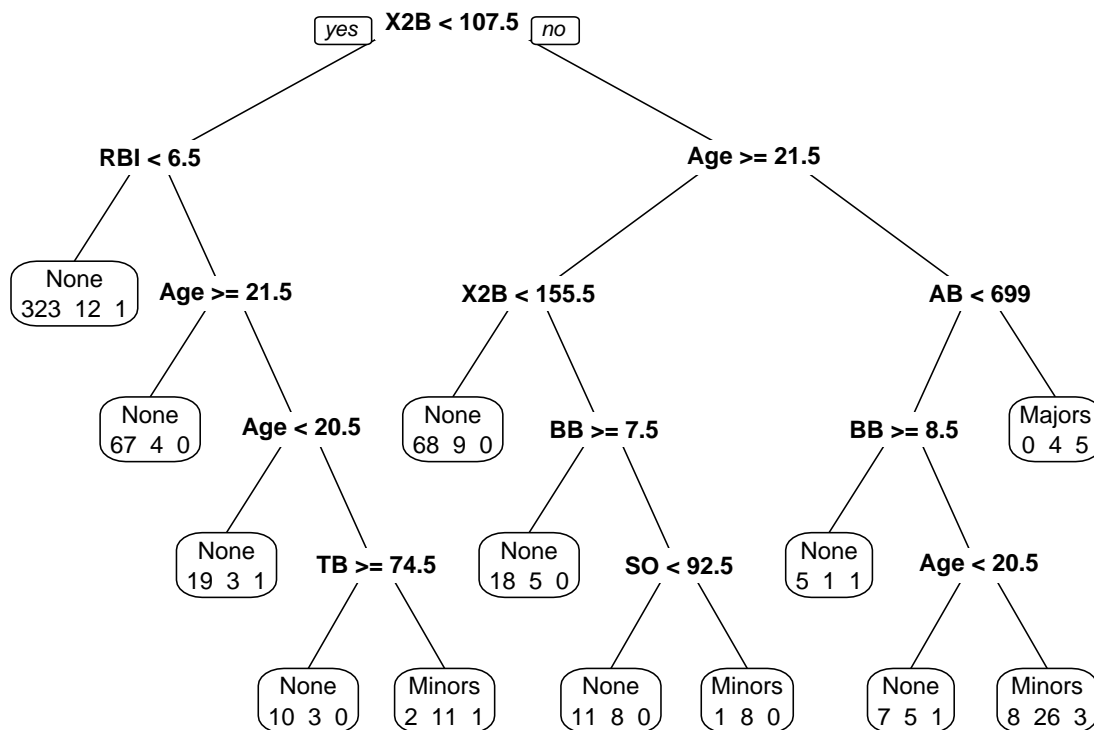
observed outcome of “None”, 96.8% were correctly predicted as none, 3.2% were incorrectly predicted as “Minors” and 0% were incorrectly predicted as “Majors”. The entries along the main diagonal indicate that the model is best at predicting players whose outcome will be “None”, and gets worse as the outcome level increases. Lastly, the proportional odds model has an overall accuracy of 91.3%.

Accuracy of PO Model for Batters

	Pred. None	Pred. Minors	Pred. Majors
Obs. None	96.8%	3.2%	0%
Obs. Minors	27.6%	69%	3.4%
Obs. Majors	22.2%	22.2%	55.6%

Random Forests

The non-parametric model is a random forest. To visualize how this method works, the outcome of a classification tree is shown below:



This graphic shows a series of yes/no questions that will result in a professional placement for every player. If a player has 107 or fewer college career doubles and 6 or fewer college career RBIs, that player will be classified as “None”. A random forest works by fitting multiple trees to different samples of the data and averaging the predictions. Here is a sample of players, their placements, and projections produced by the random forest method:

Players Actual vs Predicted Placements (Random Forest)

Name	Current.Placement	Projected.Placement
Guerrero, A.J.	None	Minors
Shimao, Tate	None	None
Matthiessen, Will	Minors	Minors
Rutschman, Adley	Majors	Majors
Reyes, Ripken	Minors	Minors
Cano, Willie	None	None
Saunders, Ty	None	None
Gibson, Hunter	None	None
Ostrander, Garret	None	None
Paganelli, Adam	None	None
Fullerton, Maxim	None	None

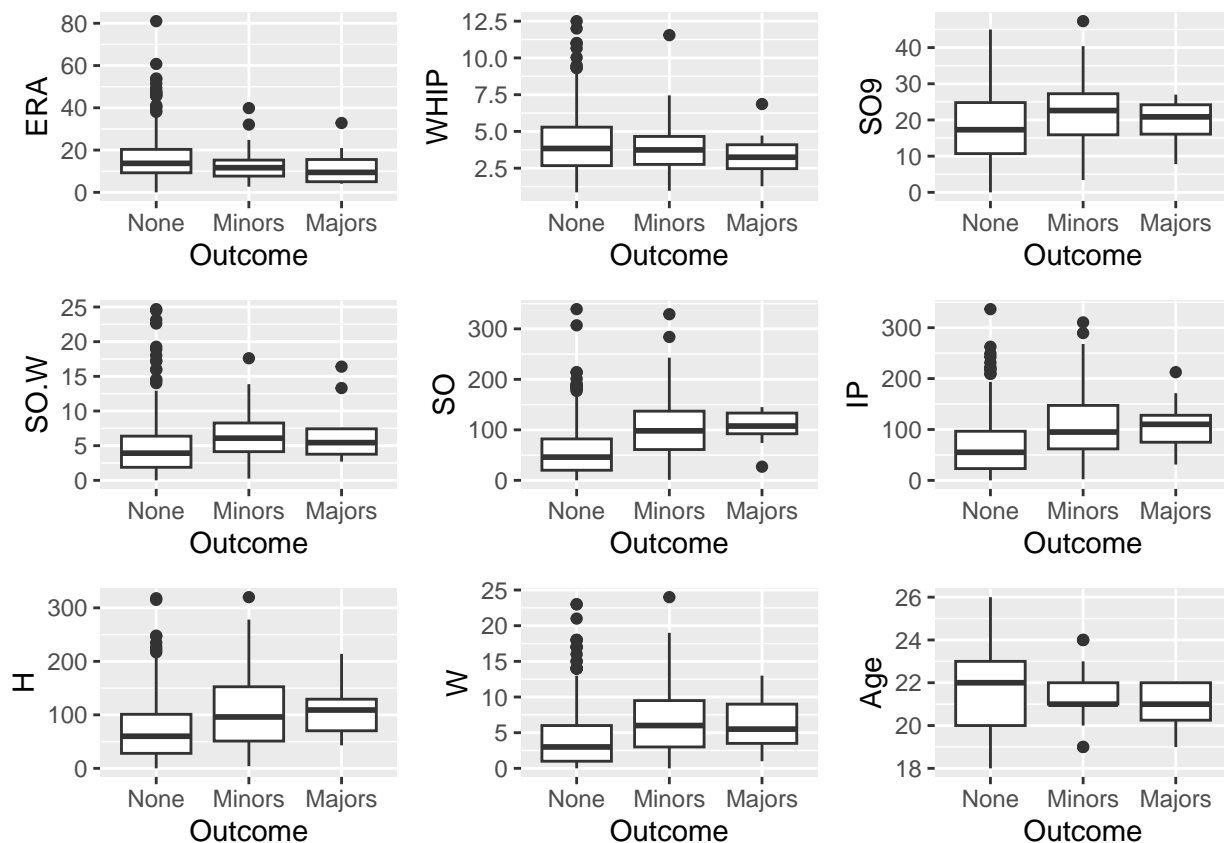
Similar to the parametric model, the table below summarizes the accuracy of this method. Random forests work significantly better than the proportional odds model, suggesting that the assumptions of the proportional odds model may not be satisfied by the data. The main diagonal shows that the random forest has some difficulty separating players who will have no professional placement from players who will be placed in the Minors, but this error is still an improvement over the proportional odds model. The random forest has an overall accuracy of 99%.

Accuracy of Random Forest for Batters

	Pred. None	Pred. Minors	Pred. Majors
Obs. None	99.8%	0.2%	0%
Obs. Minors	4.6%	95.4%	0%
Obs. Majors	0%	0%	100%

Pitchers

What follows is a repetition of the methods above applied to the data set of pitchers. ### Exploratory Analysis To see which statistics are the most related to a pitcher's professional outcome, below shows some exploratory data analysis via data visualization.



Similar to what was seen with the batters, rate-based statistics do not provide clear separation between the professional placements while counting-based statistics do. The trend relating a player's age in their final college season and placement is also seen here with the pitcher's data. With the insights from above, the two models are fit and analyzed below in the same manner as was done for the batters.

The Proportional Odds Model

A proportional odds model has been fit to the pitcher's data using the counting statistics, as indicated from the exploratory analysis above. A random selection of players, their outcomes, and projections is produced below:

Players Actual vs Predicted Placements (PO Model)

Name	Current.Placement	Projected.Placement
Haarer, Toby	None	None
Lucas, Easton	Majors	Minors
Giroux, Alex	None	None
Guarin, Marcus	None	None
Engman, Jared	None	None
Peery, Brock	None	None
Sanchez, Corey	None	None
Stumbo, Peyton	Minors	Minors
Walters, Carson	None	None

The table below shows the accuracy of this method.

Accuracy of PO Model for Pitchers

	Pred. None	Pred. Minors	Pred. Majors
Obs. None	95.8%	4.2%	0%
Obs. Minors	34.9%	65.1%	0%
Obs. Majors	14.3%	57.1%	28.6%

Again, this table shows that the proportional odds model is worse at predicting professional placements as the level of play increases. The proportional odds model for pitchers has an overall accuracy of 88.8%.

Random Forests

A random forest was fit to the pitchers data and a random sample of players, their observed placements, and projected placements is produced below.

Players Actual vs Predicted Placements (Random Forest)

Name	Current.Placement	Projected.Placement
Pinedo, Andrew	None	None
Prizina, Jake	Minors	Minors
Tongue, David	None	None
Barry, Shea	Minors	Minors
Horak, Jordan	None	None
Pilchard, Cade	None	None
Spoljaric, Garner	None	None
Mathews, Quinn	Minors	Minors

The following table summarizes the accuracy of this method:

Accuracy of Random Forests for Pitchers

	Pred. None	Pred. Minors	Pred. Majors
Obs. None	100%	0%	0%
Obs. Minors	7%	93%	0%
Obs. Majors	0%	0%	100%

Similar to what was seen for the batters, the random forest method does a better job at predicting the professional placement of pitchers. The overall accuracy of the random forest for pitchers is 99.3%.

Conclusion

This analysis demonstrates that both proportional odds models and random forests can provide useful predictions of professional placement for NCAA baseball players. Counting-based statistics (such as home

runs, RBIs, innings pitched, and strikeouts) and player age in the final NCAA season consistently emerged as strong indicators of future professional outcomes.

Random forests achieved near-perfect accuracy on the training data, substantially outperforming the proportional odds model. However, such performance is almost certainly the result of **overfitting**, meaning the model is capturing noise in the training data rather than patterns that will generalize to new players. Proportional odds models, although less accurate, are more interpretable and provide confidence in identifying which statistics meaningfully influence player outcomes.

For practical applications such as recruitment, these results suggest caution. Random forests may be useful as an exploratory tool or for screening players, but they must be validated with cross-validation or out-of-sample testing before being relied upon for decision-making. The proportional odds model offers complementary value by explaining the drivers of placement in a statistically rigorous way.

Future work should focus on validating these models on unseen data, incorporating longitudinal statistics to capture player development, and integrating qualitative metrics (such as scouting reports). This would improve both predictive accuracy and real-world applicability.

This analysis suggests certain key statistics as holding significant power in projecting professional placement of NCAA Baseball players. Ultimately, statistical modeling offers a valuable complement to traditional scouting, but its insights are most reliable when balanced with domain expertise and rigorous validation.